

# New York Times Articles & Comments (2020)

A data-driven analysis of reader interaction patterns across New York Times content from 2020.

by Akash K P, Luv Aggarwal, Rachit Mehta, Vivek Rajput





# Problem Statements

The New York Times receives millions of comments annually. Understanding and predicting reader engagement can improve content strategy and user experience.



## Comment Prediction

Predict the number of comments an article will receive



## Recommendation Count

Predict how many recommendations a comment will get



## Editor's Pick Classification

Predict which comments editors will select as Times Picks



## Headline Generation

Generate headlines using LSTM neural networks





# Task 1: Predicting No. of Comments Using Machine Learning

**Goal:** Use metadata from NYT articles to predict reader engagement (comment count).

**Pipeline:**


- Cleaned & engineered text features
- Built a scalable ML pipeline
- Models used: Linear Regression, Neural Network, Random Forest, XGBoost

**Feature Engineering (novelty):**

- Extracted word counts from the headline and abstract
- Created interaction term: word count × headline word count
- One-hot encoded the section categorical column
- Standard scaled numerical features for model compatibility

**Results After Feature engineering:**

Model	R <sup>2</sup> Score
Linear Regression	0.3361
Neural Network	0.3794
Random Forest	0.28
XGBoost	0.3710



**Target Variable Transformation:**

- The target variable (number of comments) had a **right-skewed distribution**
- Applied **logarithmic transformation**:  $y = \log(1 + \text{number of comments})$
- Helped normalise the target for better model learning, Reduced the impact of extreme values (outliers)
- Resulted in a **significant improvement in R<sup>2</sup> score** and overall model performance

Model	R <sup>2</sup> Score
Linear Regression	0.4183
Neural Network	0.4712
Random Forest	0.4277
XGBoost	0.4713



# Task 1: Predicting No. of Comments Using Machine Learning

## TF-IDF Vectorization :

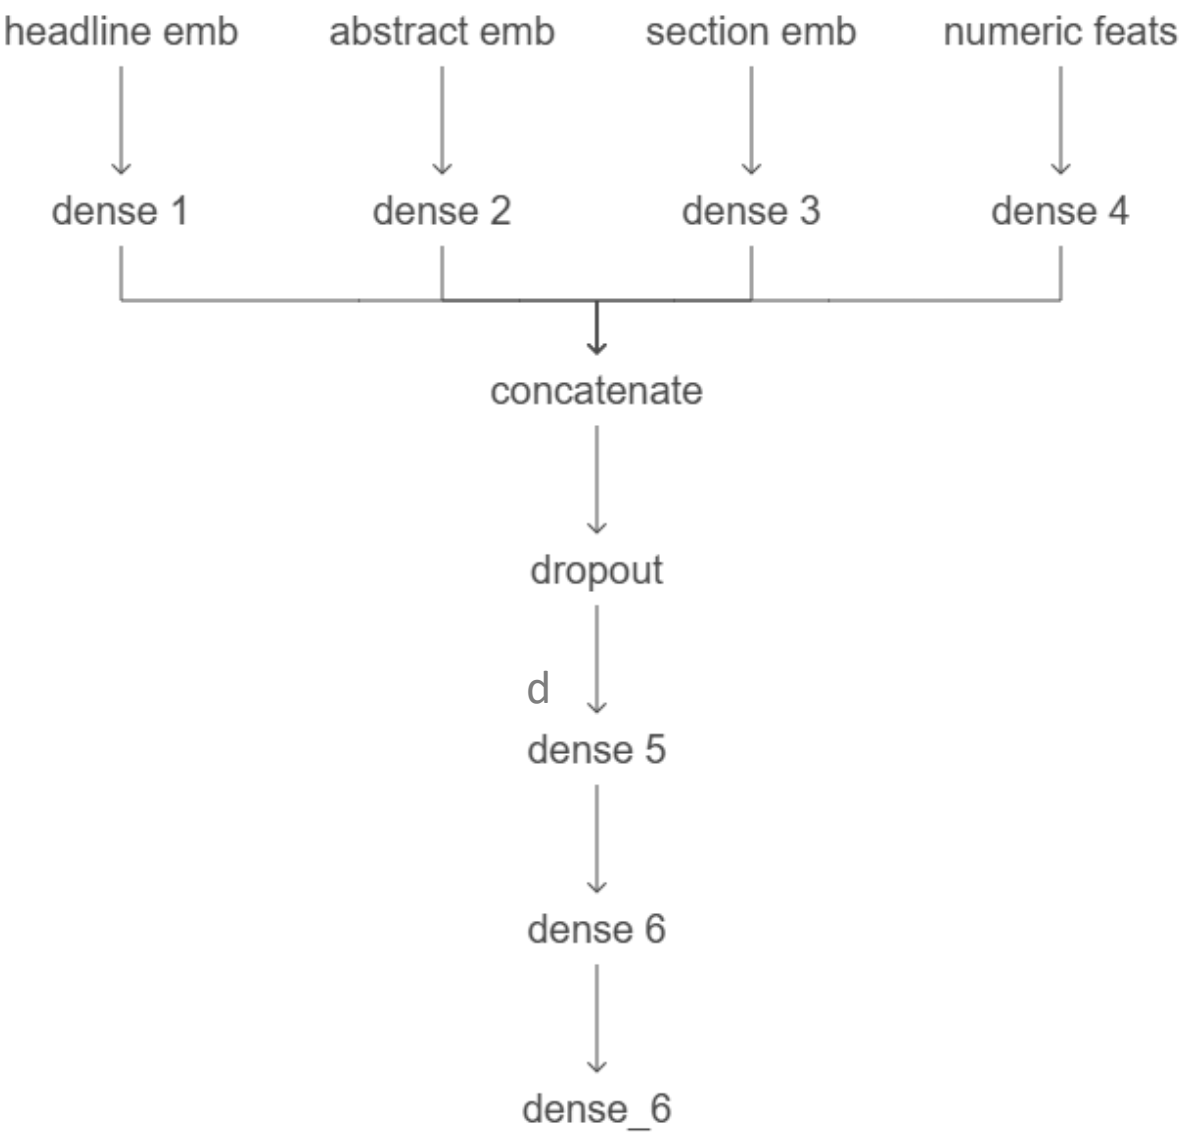
- Applied **TF-IDF (Term Frequency–Inverse Document Frequency)** to convert text data into a numerical format
- Focused on headline and/or abstract columns for feature extraction
- Captured the importance of words relative to the entire dataset (not just raw counts)
- Reduced the influence of common words and highlighted meaningful terms
- Enabled models to understand text features like article titles and summaries
- Improved prediction performance by turning text into structured input

Model	R <sup>2</sup> Score
Linear Regression	0.4148
Neural Network	0.4722
Random Forest	0.4237
XGBoost	0.4658

## Architecture 1 using Fusion layers for embedding:

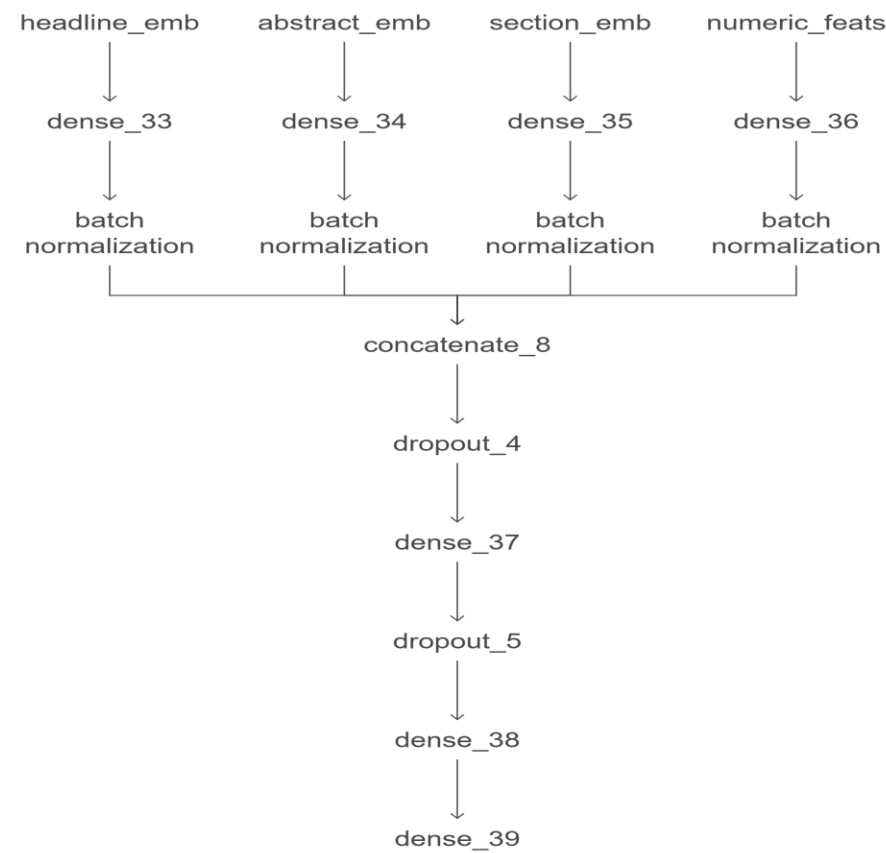
- Used **BERT** Model sentence-transformers All mini LM L6 v2 to generate embeddings
- Then used fusion dense Neural Network layers and applied dropout after concatenation.
- R<sup>2</sup> Score: 0.49**

## Neural Network Architecture



# Task 1: Predicting No. of Comments Using Machine Learning

Neural Network Layer Connections

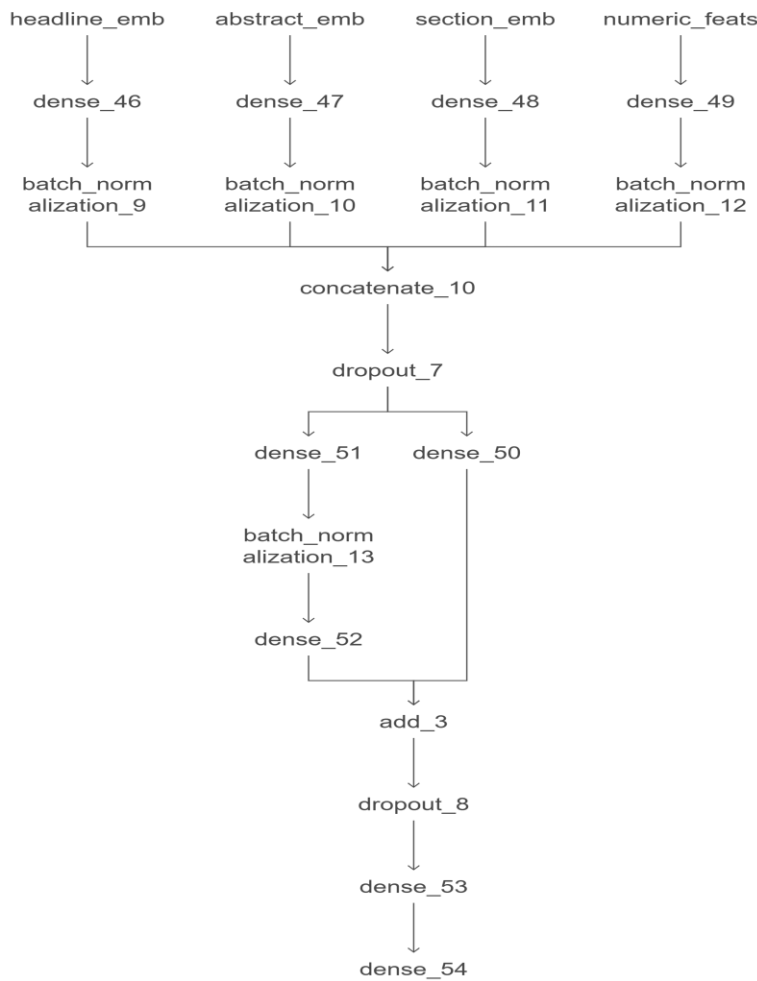


Made with Napkin

## Architecture 2 using Fusion layers for embedding:

- Used Batch Normalisation.
- **R<sup>2</sup> Score: 0.50** with more numerical features.

Neural Network Layer Connections



Made with Napkin

## Architecture 3 using Fusion layers for embedding:

- Used Residual Learning and Skip Connection.
- **R<sup>2</sup> Score: 0.5325**

# Task 1: Predicting No. of Comments Using Machine Learning

## BERT Embeddings:

- Used **pre-trained BERT embeddings** to convert text (e.g., headlines) into dense vectors
- Captured **contextual meaning** of words based on their position and usage in a sentence
- Provided richer, more accurate text representation compared to traditional methods
- Used as input for **deep learning models**, boosting prediction accuracy

Model	R <sup>2</sup> Score
Linear Regression	-3.01 e18
Neural Network	-0.0637
Random Forest	0.5417
XGBoost	0.5474

## BERT with Feature Engineering:

- Extracted **BERT embeddings** from textual data like headlines and/or abstracts
- Engineered a custom **interaction term**: word count × headline word count
- Combined **BERT vectors** with structured features (e.g., interaction term, word counts) into a unified input
- Enabled models to learn from both **semantic text context** and **quantitative patterns**
- Resulted in improved model performance by leveraging both deep language understanding and structured signals

Model	R <sup>2</sup> Score
Linear Regression	-1.04 e18
Neural Network	-0.046
Random Forest	0.5419
XGBoost	0.5499

# Task 1: Predicting No. of Comments Using Machine Learning


## Tuned Bert Embeddings:

### Using Raw BERT / DistilBERT (AutoModel from Hugging Face)

- Outputs last hidden states (sequence of token embeddings)
- Optionally includes the [CLS] token representation
- Provides **raw contextualized word representations**
- We get **token-level embeddings** for each input token

### Flexibility and Control:

- We **fine-tune** the transformer model
- Multiple strategies to extract sentence-level representations:
  - **Mean Pooling:** Average across token embeddings
  - **[CLS] Token:** Use embedding of the [CLS] token
  - **Attention Pooling:** Learnable attention weights to pool token representations
- Gives **more control** over how final features are formed
- R<sup>2</sup> Score: 0.5589**





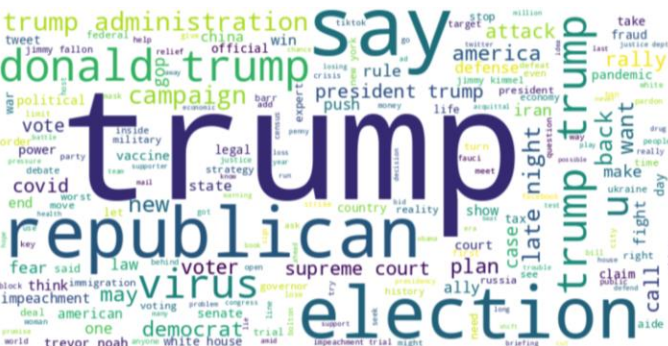
## Task 2: Unsupervised Topic Discovery in NY Times Headline (2020)

## Goal: To identify hidden themes in NY Times headlines from 2020 using **clustering** and **topic modeling**

## Elbow Method

- The Elbow Method helps determine the optimal number of clusters by plotting inertia against different K values.
- Inertia decreases as the number of clusters increases, but the rate of decrease slows after a certain point.
- The “elbow” is observed at **K = 5**, indicating it as the optimal number of clusters.



## K-Means Clustering:



## Cluster 1



## Cluster 2



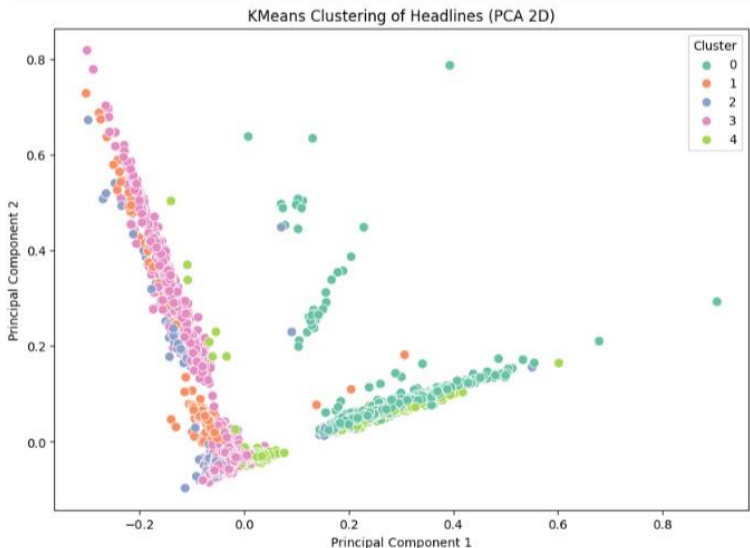
### Cluster 3



## Cluster 4




## Cluster 5



# Task 2: Unsupervised Topic Discovery in NY Times Headline (2020)

## LDA Clustering:



## Comparing K-Means Clustering vs LDA Topic Modelling:

Feature	KMeans Clustering	LDA Topic Modeling
Clustering Type	Hard clustering (each doc → 1 cluster)	Soft clustering (each doc → mix of topics)
Vectorization Used	TF-IDF (Term Frequency–Inverse Document Frequency)	Count Vectorizer (word frequency)

		headline \				
10156		Picasso Mural Torn From Building After Years of Dispute				
15209		Nail Salons, Lifeline for Immigrants, Have Lost Half Their Business				
5267		Kicked Out of China				
11723		Tiwa Savage, Queen of Afrobeats, Makes a New Start				
16039		2020: The Year in Sports When Everybody Lost				
	cluster	Topic_0	Topic_1	Topic_2	Topic_3	Topic_4
10156	3	0.733332	0.066667	0.066667	0.066667	0.066667
15209	3	0.050379	0.051229	0.549085	0.299089	0.050218
5267	3	0.100000	0.100000	0.100993	0.101753	0.597254
11723	3	0.302284	0.040579	0.361586	0.040615	0.254936
16039	3	0.798221	0.050450	0.051099	0.050230	0.050000



# Task 3: Deep Learning-Based Text Generation on NYT Dataset



**Goal:** Generate news headlines using a Long Short-Term Memory (LSTM) neural network.

## Embedding:


- Used **Keras Tokenizer** to convert text into integer sequences (word index).
- Passed tokenized input through an **Embedding layer** which maps each word index to a dense vector, capturing semantic relationships between words.

## Model Training:

Input → Embedding Layer → LSTM Layer → Dense Layer (ReLU) → Dense Layer (Softmax)

## Semantic Safety Filtering:

- Used a zero-shot classification pipeline from Hugging Face Transformers.
- Each generated headline was checked against a list of sensitive topics: e.g. violence, tragedy, politics, hate speech
- e.g. **“tech stock look under Russia stays in protest protesters can’t stop.”** (without safe mode)



## NYT Headline Generator

Generate news headlines using your custom-trained LSTM model. You can optionally filter out unsafe topics.

Enter a seed phrase:

tech stocks

Number of words to generate

10

520

Temperature

1.00

0.501.50

☐ Generate safe headline only

Generate Headline

Generated Headline:

tech stocks look under russia stays in protest protesters can't stop



# Task 3: Deep Learning-Based Text Generation on NYT Dataset



**Goal:** Generate news headlines using a Long Short-Term Memory (LSTM) neural network.


Headline Generation:

- Input a seed phrase.
- Predicts one word at a time using previously generated words.
- Used temperature sampling:
  - Low temperature is more focused & repetitive
  - High temperature is more diverse & creative
- Controls for Generation: Number of Words, Safe Mode

Example:

User Seed: tech stocks

Generated Headlines: “***tech stocks have it hard to go back to power trump***” (with safe mode)



## NYT Headline Generator

Generate news headlines using your custom-trained LSTM model. You can optionally filter out unsafe topics.

Enter a seed phrase:

tech stocks

Number of words to generate

10

520

Temperature

1.00

0.501.50

☒ Generate safe headline only

Generate Headline

Generated Headline:

tech stocks have it hard to go back to power trump

# Task 4: Editor’s Pick Prediction

**Goal:** Predict whether a New York Times comment will be selected as an **Editor’s Pick**.

## Initial Experiments

### 50,000 Points Dataset (XG Boost Classifier Model)

- Balanced Performance with Class 0 Leading: Achieved 92% accuracy with solid performance for Class 0.
- Class 0 Precision: 0.94, Recall: 0.95, F1-score: 0.95
- Class 1 Precision: 0.85, Recall: 0.82, F1-score: 0.84
- Strong Weighted Performance: Weighted F1-score: 0.92, indicating a balanced overall model despite class imbalance.

### Full Dataset (Times Pick Data)

- High Overall Accuracy with Class Imbalance: Achieved 96.74% accuracy, though impacted by class imbalance.
- Class 0 Precision: 0.9889, Recall: 0.9780, F1-score: 0.9834
- Class 1 Precision: 0.0257, Recall: 0.0502, F1-score: 0.0340
- Imbalance Challenge: Model is heavily biased toward Class 0, requiring strategies to improve Class 1 detection.

## •BERT-Based Modelling

### •BERT + metadata model

- Significant improvement in recall for class 1: 85.91%
- Precision: 0.1627, Recall: 0.8591,
- F1-score: 0.2736
- Maintained strong overall performance (accuracy: 94.78%)
- Demonstrated that combining contextual text modelling with metadata significantly improves minority class performance

	precision	recall	F1-score
0	0.94	0.95	0.95
1	0.85	0.82	0.84
Accuracy		0.92	
Macro avg	0.9		
Weighted avg	0.92		

	precision	recall	F1-score
0	0.9889	0.9780	0.9834
1	0.0257	0.0502	0.0340
Accuracy			0.9674

	precision	recall	F1-score
0	0.9983	0.9488	0.9729
1	0.1627	0.8591	0.2736
Accuracy		0.9478	
Macro avg	0.5805		
Weighted avg	0.9887		

•Accuracy : 0.9478

# Task 4: Editor’s Pick Prediction


**Goal:** Predict whether a New York Times comment will be selected as an **Editor’s Pick**.

## Threshold Tuning Phase

- Performed precision-recall vs threshold analysis
  - Adjusted the classification threshold (0.7) to better balance recall and precision
  - Retrained the model using the optimized threshold
- Final model performance:
  - Class 1 (Editor’s Picks):
    - Precision: 0.2590
    - Recall: 0.7523
    - F1-score: 0.3853 (highest so far)
  - Class 0 retained high performance (F1-score: 0.9860)
  - Overall accuracy: 97.25%
  - Macro F1-score: 0.6856
  - Weighted F1-score: 0.9791

## Techniques and Key Insights

- Addressed a highly imbalanced classification problem
- Leveraged BERT for advanced contextual understanding of text
- Augmented BERT with metadata features for better model grounding
- Tuned classification thresholds using precision-recall trade-off analysis
- Iteratively improved model performance on both majority and minority classes
- Achieved a well-balanced, production-ready model with high accuracy and strong recall for editor’s picks



	precision	recall	F1-score
0	0.9971	0.9751	0.9860
1	0.2590	0.7523	0.3853
Accuracy			0.9725
Macro avg	0.6280	0.8637	0.6856
Weighted avg	0.9886	0.9725	0.9791

•Accuracy : 0.9725



# Future Work

- Ensemble Methods: Apply ensemble techniques (e.g., stacking, boosting) to combine multiple models for more.
- Model Improvements: Replace or augment LSTM models with Transformer-based architectures (e.g., GPT, T5) to generate higher-quality, more coherent headlines.
- Model the full thread or parent-child structure of comments using techniques like graph neural networks (GNNs)